



Convolutional networks

David Tonderski (davton)



(a) Training information for network 1. The training took under 2 minutes.



(b) Training information for network 2. The training took under 2 minutes.

Figure 1: Training information for the two networks.

After training, network 1 has a training accuracy of about 98.43%, a validation accuracy of 97.85%, and a test accuracy of 98.18%, while network 2 has a training accuracy of about 99.01%, a validation accuracy of 98.6%, and a test accuracy of 98.88%. As expected, the deeper second network performs slightly better. Interestingly, the training times using the GPU were quite similar, which I did not expect from the description (2 hours vs 6 hours).